



An interpretable Dictionary Learning approach for Diagnostic Hypotheses from 3D Chest

An undergraduate project report submitted to the

Department of Electrical and Information Engineering
Faculty of Engineering
University of Ruhuna
Sri Lanka

in partial fulfillment of the requirements for the

**Degree of the Bachelor of the Science of Engineering
Honours**

by

K.R.C.S. Kekulawala - EG/2021/4609
M.D.Y.B. Kularathne - EG/2021/4620
B.G.P.U. Thilakshana - EG/2021/4831
B.P.G.L. Thejana - EG/2021/4828

11th May 2026

.....
Dr. Kaveen Liyanage
(Supervisor)

.....
Mrs. Sithara Mahagama
(Co-Supervisor)

Abstract

Multi-class abnormality detection from 3D chest Computed Tomography requires clinically interpretable methods to support radiological decision-making. Current state-of-the-art (SotA) models often rely on post-hoc explainability techniques that provide approximate localizations and lack direct traceability to the models decisions. Furthermore, SotA models rarely support simultaneous multi-class classification and grounded localization, leading to a fragmented approach to diagnosis. This study aims to address these challenges through two contributions: First, we introduce an inherently interpretable sparse dictionary learning framework in which detected abnormalities, such as Pulmonary Nodules, Ground-Glass Opacity and Consolidation, can be back-projected to the original voxel space, allowing diagnostic decisions to be traced back to relevant anatomical structures. Second, we develop a unified evaluation framework to benchmark SotA models across three axes: (i) classification accuracy, (ii) segmentation quality, and (iii) explainability, using standardized metrics and a publically available 3D chest CT dataset with expertly annotated abnormality localizations. Our method is benchmarked against SotA models using this framework, demonstrating competitive classification and segmentation performance while providing improved interpretability compared to deep learning baselines.

Acknowledgments

We would like to express our sincere gratitude to Dr. Kaveen Liyanage (Supervisor), Mrs. Sithara Mahagama (Co-Supervisor), and Dr. Subhodini Indrapala (Consultant Radiologist at Teaching Hospital, Karapitiya), for their invaluable guidance, support, and constructive feedback throughout the course of this project. Their expertise, mentorship, and encouragement have been instrumental in shaping the direction and quality of this research.

We are grateful to the Department of Electrical and Information Engineering for providing access to computing resources and research facilities that were essential for conducting this project. The department’s commitment to fostering an environment conducive to research and innovation has been greatly appreciated.

We would like to acknowledge the creators and maintainers of the publicly available datasets used in this research, particularly the CT-RATE and RexGrounding-CT datasets and other medical imaging resources that have made this work possible. These datasets have been crucial for developing and validating the proposed methods.

We extend our thanks to the developers of the open-source frameworks and libraries used in this project, including PyTorch, scikit-image, and other tools that significantly contributed to the implementation of this research.

Contents

Abstract	ii
Acknowledgements	iii
Contents	v
List of Figures	vi
List of Tables	vii
Acronyms	viii
1 Introduction	ix
1.1 Background	ix
1.2 Problem Statement	ix
1.3 Objectives and Scope	x
1.3.1 Objectives	x
1.3.2 Scope	x
1.4 Literature Review	x
1.4.1 Classification Architectures	xi
1.4.2 Segmentation Architectures	xi
1.4.3 Unified Architectures	xii
1.4.4 Dictionary Learning Approaches	xii
1.4.5 Explainability Approaches	xii
1.5 Report Outline	1
2 Background	2
2.1 Dictionary Learning	2
2.2 Radiological Patterns in Lung CT Imaging	3
2.2.1 CT Imaging	3
2.2.2 Ground Glass Opacity (GGO) and Consolidation	3
2.2.3 Pulmonary Nodules	3
2.3 Vision Transformers (ViTs)	4
3 Methodology	5
3.1 Data Acquisition and Preprocessing	5
3.2 Implementation of the LC-KSVD Approach	6
3.3 Abnormality Classification	6
3.4 Back-Projecting Class-Discriminative Atoms	6
3.5 Deriving Explainability through Segmentation Masks	7

<i>CONTENTS</i>	v
4 Experimentation and Evaluation	8
4.1 Grounded Localization Evaluation Results	8
4.1.1 Overall Performance Summary	8
4.1.2 Per-Class Performance	9
4.1.3 Distribution Analysis	9
4.1.4 Critical Finding: Focal vs. Non-Focal Lesion Bias	9
4.1.5 Score Distribution and Failure Modes	9
4.1.6 Conclusions from Localization Evaluation	10
5 Discussion	13
A Appendix 1	14
B Appendix 2	15
APPENDIX	15
Bibliography	16

List of Figures

4.1	Per-class localization performance across three key metrics. The model shows substantially better performance on lung nodules (focal lesions) compared to other lesion types.	9
4.2	Histogram of Dice scores for each lesion class. Red and blue dashed lines indicate mean and median respectively. Zero-Dice percentages shown in subtitle. Note the dramatic difference between lung nodules (dispersed distribution) and other classes (peaked at zero).	10
4.3	Critical performance gap between focal (lung nodules) and non-focal (diffuse) lesions. Left panel shows direct metric comparison. Right panel visualizes the performance multiplier gap: the model achieves 18.1× higher Dice on nodules versus non-focal lesions.	11
4.4	Cumulative Distribution Function of Dice scores. Shows the proportion of cases achieving each Dice threshold. Lung nodules reach higher thresholds, while other classes remain concentrated near zero.	11
4.5	Stacked bar chart showing the fraction of cases in each performance category. Consolidation and atelectasis show greater than 98% zero-Dice cases (complete failure to localize). Lung opacity shows 74% zero-Dice. Only lung nodules show meaningful detection diversity. . .	12

List of Tables

4.1	Localization Performance by Lesion Class	9
4.2	Morphology-Based Performance Analysis	9

Acronyms

AC - Alternating Current
DC - Direct Current
LV - Low Voltage
MV - Medium Voltage

Chapter 1

Introduction

Developing a unified architectural pipeline for multi-class abnormality classification and grounded localization in 3D Computed Tomography (CT) addresses two of the most critical imperatives in clinical AI adoption: diagnostic accuracy and explainability. Specifically, utilizing dense pixel-level segmentation as the mechanism for localization forces the architecture to provide voxel-level visual evidence for its diagnostic claims. This ensures that the model is held accountable and its reasoning aligns with clinical diagnosis.

1.1 Background

The transition from 2D projectional radiography to volumetric CT represents a paradigm shift in medical image analysis. For Lung disease diagnosis, Chest X-rays (CXR) are commonly used as the first imaging test due to its low cost and low radiation exposure. However, studies show that their capability to detect abnormalities is limited when compared to 3D CT. [1].

Furthermore, CXR has historically benefited from the availability of massive, radiologist-annotated datasets, while the analysis of chest CT volumes has been severely constrained by a lack of such data. However, the recent introduction of large-scale, text-paired datasets has facilitated the development of state-of-the-art

(SotA) vision-language foundation models tailored specifically for volumetric imaging. [2] [3].

Beyond deep learning approaches, Dictionary Learning algorithms, particularly, Label Consistent KSVD (LC-KSVD) [4], Label Embedded Dictionary Learning (LEDL) and Fischer Discriminant Dictionary Learning (FDDL) offer viable approaches for interpretable CT analysis. All three techniques incorporate label information during the dictionary training process, and ensures that the sparse representation of the data is highly discriminative, making them well-suited for multi-abnormality tasks.

1.2 Problem Statement

The clinical utility of Deep Learning-based approaches is hindered by a lack of interpretability. While tools like Grad-CAM provide post-hoc explanations to model decisions, these are often insufficient for medical professionals who require a granular understanding of the model’s decision making process. In high stakes environments, a model that produces a diagnosis without transparency in its reasoning is difficult

to integrate into standard workflows, as it lacks the accountability necessary for legal and ethical compliance.

1.3 Objectives and Scope

The primary goal of this project is to address the limitations of existing Deep Learning based approaches for Lung abnormality classification and localization by developing an inherently interpretable dictionary learning framework. This approach aims to enable transparent, anatomically grounded decision making, where each classification outcome can be traced back to specific pathological patterns. This aligns with clinical reasoning and facilitates validation and adoption in radiological workflows. The specific objectives are as follows:

1.3.1 Objectives

- Implement a Dictionary Learning approach to learn normal and pathological dictionary atoms using a publicly available annotated 3D chest CT dataset.
- Map dictionary atoms back to anatomical features and visualize their spatial contribution maps.
- Benchmark model performance against state-of-the-art classification and segmentation models on the same dataset.
- Evaluate abnormality localization by the model on a real-world dataset with expert radiologist review.

1.3.2 Scope

The scope of this study focuses on the development and evaluation of a dictionary learning based approach for classification and grounded localization of lung abnormalities using volumetric 3D CT data, with the following limitations:

- This study is limited to 3 thoracic abnormalities: Pulmonary Nodules, Ground-Glass Opacities, and Consolidations, based on expert radiologist recommendation and dataset availability.
- This study will utilize a single publicly available dataset for training and evaluation, which may limit generalizability to other populations and scanner types.
- The proposed approach will be limited only to linear methods, and will not explore non-linear or hybrid models that provide better performance for less explainability.

1.4 Literature Review

The landscape of 3D CT analysis is currently divided into three primary architectural paradigms:

- Classification-based models that prioritize diagnostic labels

- Segmentation-based models that prioritize boundary precision
- Unified architectures designed for grounded localization

SotA models are predominantly specialized in either classification or segmentation. However, recent systems, particularly in 2025 and 2026, are moving toward multi-task learning or foundation models that can perform both diagnosis and voxel-level segmentation simultaneously.

Table 1 summarizes the key models in each category, their primary task orientation, and their performance on standard benchmarks.

1.4.1 Classification Architectures

Classification-based models for 3D CT have shifted away from 3D Convolutional Neural Networks (CNNs) such as 3D U-Net, 3D ResNet and DenseNet pipelines, toward Vision Transformers (ViTs) and self-supervised encoders. The benchmark for this paradigm is the CT-RATE dataset, which introduced the first public large-scale 3D medical imaging dataset with paired textual reports, enabling the development of Contrastive Language-Image Pretraining for Computed Tomography (CT-CLIP).

CT-CLIP is a VLFM designed for 3D chest CT analysis. It extends the CLIP framework to volumetric medical imaging by jointly learning representations from CT scans and their corresponding radiology reports. [5, 6] Through contrastive learning, the model aligns image features with textual descriptions, enabling tasks such as zero-shot classification, cross-modal retrieval, and report-guided abnormality localization. [5] The model is particularly effective in learning semantic relationships between radiological findings and volumetric imaging patterns such as GGO, consolidation, fibrosis, and pulmonary nodules. [5, 7]

Recent benchmarks such as Merlin and DINOv3 have further challenged the necessity of domain-specific pretraining. Merlin is a VLFM exclusively pre-trained on abdominal CT scans. Despite its training, Merlin has consistently demonstrated exceptional zero-shot performance on chest CT classification tasks, even when it was restricted to a linear probe, significantly outperforming models like CT-CLIP, M3FM, and CT-FM that were explicitly trained on thoracic data, by an average of 12.3% in Area Under the Receiver Operating Characteristic Curve (AUROC). Specifically, in cross-domain retrieval tasks, Merlin surpassed CT-CLIP by 24.7%.

DINOv3, a self-supervised ViT pretrained on over 1.7 billion natural images, outperform specialized medical models like CT-Net and CT-CLIP across all metrics on several classification tasks due to its ability to capture fine-grained textural features.

1.4.2 Segmentation Architectures

The prevailing SotA for supervised volumetric segmentation remains variations of the 3D U-Net, such as nnU-Net, which utilizes an adaptive framework to optimize preprocessing and architecture for specific datasets. However, the field has recently shifted toward foundation models capable of universal segmentation such as SegVol and SAM-Med3D. The SegVol model is the first 3D foundation model specifically designed for universal volumetric medical image segmentation. Similarly, SAM-Med3D adapts the Segment Anything Model (SAM) for 3D medical data.

Recent evaluations published at the Machine Learning for Healthcare (MLHC) Conference in 2025, which specifically utilized the ReXGroundingCT dataset, revealed that the current SotA text-prompted segmentation models that demonstrate near-perfect performance in general computer vision tasks struggle with the nuanced clinical descriptions and ill-defined boundaries of chest CT abnormalities.

1.4.3 Unified Architectures

Recent models have been designed to integrate abnormality detection, segmentation, and multimodal chain-of-thought reasoning, enabling pixel-level localization, structured report generation, and physician-like diagnostic inference in a single framework. BiomedParse, and its cloud-deployed version MedImageParse, developed by Microsoft, are foundation models designed to parse 3D medical images. Other notable models include Citrus-V, MTMed3D and MedVista3D, where each architecture incorporates a multi-task learning (MTL) framework.

Table 1

1.4.4 Dictionary Learning Approaches

Applications of sparse dictionary learning in medical imaging have demonstrated significant success across diverse diagnostic tasks. In histopathological analysis, Discriminative Feature oriented Dictionary Learning (DFDL) has shown the ability to learn class specific dictionaries that achieve high classification accuracy even with limited training data a critical advantage when generous training samples are not available[19]. In neuroimaging, multi feature kernel supervised dictionary learning approaches have demonstrated superior performance in distinguishing Alzheimer’s disease patients from cognitively unimpaired individuals, achieving 98.18% classification accuracy while maintaining fast testing times[20].

Most relevant to our work, dictionary learning has been successfully applied to lung disease diagnosis from CT images. One approach utilized both frozen dictionary learning and label consistent LC-KSVD to classify COVID 19 CT scans, achieving 89% accuracy on the CC-CCII dataset[21]. The framework’s key advantage lies in its interpretability, where the sparse coding output can be directly mapped back to the original input signal domain. This work demonstrates that sparse representation methods offer an effective balance between performance and explainability for lung disease classification tasks, providing clear feature representations while maintaining competitive accuracy. Similarly, sparse representation methods have achieved 95.4% accuracy in classifying diffuse lung disease patterns on HRCT images, effectively handling the high variety and complexity of patterns[22].

1.4.5 Explainability Approaches

Grad-CAM is one of the most widely used post-hoc explanation methods for CNN-based classifiers [8]. It generates class-discriminative heatmaps by using gradients flowing into the final convolutional layer, allowing the user to identify image regions that influenced a prediction [8]. Its main limitation is that it produces low-resolution saliency maps and does not directly provide precise lesion boundaries or voxel-level localization, which is a major drawback for 3D CT interpretation [8?].

Seg-Grad-CAM extends this idea by adapting Grad-CAM to semantic segmentation outputs [?]. Instead of only highlighting regions that are important for image-level classification, it produces pixel-level relevance heatmaps for segmentation predictions [?]. This is particularly useful for volumetric medical imaging tasks where the target abnormality occupies a sparse region within a large 3D scan.

HiRes-CAM addresses one of the key weaknesses of Grad-CAM by producing higher-resolution and more faithful localization maps [?]. This is important in chest CT because many abnormalities, such as small nodules, consolidations, or subtle interstitial changes, may be missed by coarse heatmaps. HiRes-CAM aims to preserve fine spatial detail, making it more suitable for medical imaging applications where precise localization is required [?]. However, like other gradient-based methods, it remains a post-hoc explanation technique and does not change the internal decision process of the model.

1.5 Report Outline

This report is organized into five main chapters.

Chapter 1 introduces the research problem, background, objectives, scope, and literature review related to multi-class abnormality classification and grounded localization in 3D chest CT imaging. It also discusses existing deep learning and dictionary learning approaches together with their limitations in explainability.

Chapter 2 presents the theoretical background required for this study, including sparse representation, dictionary learning, K-SVD, LC-KSVD, frozen dictionary learning, and Vision Transformers used in medical imaging. In addition, important radiological abnormalities such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules are discussed.

Chapter 3 describes the proposed methodology of the project. This includes dataset acquisition, preprocessing steps, patch extraction, implementation of the LC-KSVD framework, abnormality classification pipeline, and the proposed back-projection mechanism used for interpretable localization and explainability generation.

Chapter 4 explains the experimentation and evaluation framework used in this research. It presents the selected baseline models, evaluation metrics, and the three evaluation axes: classification accuracy, grounded localization quality, and explainability. The chapter also summarizes the current progress of the project and planned evaluations.

Finally, Chapter 5 discusses the overall findings, contributions, limitations, and future improvements of the proposed interpretable dictionary learning framework for 3D chest CT abnormality analysis.

Chapter 2

Background

2.1 Dictionary Learning

Dictionary learning is a signal modeling technique where a CT image patch is expressed as a linear combination of a small number of basis elements called dictionary atoms. [9] Instead of using all available features, the signal is approximated as:

$$y = Dx$$

where D represents the dictionary and x is a sparse coefficient vector containing mostly zero values. This produces a compact and interpretable representation. [10]

The dictionary D is typically learned using K-SVD (K-Singular Value Decomposition), which solves the following optimization problem:

$$\min_{D, X} \|P - DX\|_F^2 \quad \text{s.t.} \quad \forall i, \|x_i\|_0 \leq T$$

where P represents the training data, X is the sparse coefficient matrix, and T controls the sparsity level. K-SVD alternates between sparse coding and dictionary update steps to minimize reconstruction error. However, it does not incorporate class labels, which limits its discriminative ability for classification tasks. [9]

It jointly learns a dictionary, sparse codes, and a classifier so that the learned representations are not only good for reconstruction but also discriminative for classification through label consistency constraints. However, LC-KSVD still learns a single shared dictionary for all classes.

To address this limitation, we adopt a two-stage dictionary learning strategy inspired by constrained dictionary learning approaches. [11] In the first stage, a dictionary is learned from normal CT scans to capture healthy anatomical structures. In the second stage, this dictionary is frozen and kept fixed, while additional atoms are learned only from abnormal scans. This approach can be interpreted as:

- preserving D_{normal} learned from healthy data
- learning D_{abnormal} independently for disease patterns

This separation ensures that normal anatomical structures are not overwritten, while abnormal features are explicitly isolated, improving interpretability and stability.

2.2 Radiological Patterns in Lung CT Imaging

2.2.1 CT Imaging

Computed Tomography (CT) imaging is widely used in lung disease diagnosis due to its ability to generate detailed cross-sectional images of lung structures. [12] CT scans provide high-resolution visualization of abnormalities such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules, enabling accurate assessment and early detection of lung diseases. [13]

In modern medical imaging systems, CT imaging also supports computer-aided diagnosis and deep learning-based analysis for automated detection, segmentation, and classification of lung abnormalities. [14] These capabilities improve diagnostic efficiency and assist radiologists in clinical decision-making.

2.2.2 Ground Glass Opacity (GGO) and Consolidation

Ground Glass Opacity (GGO) and consolidation are two important radiological patterns used in lung CT diagnosis. [13]

Ground Glass Opacity (GGO): [15]

- Appears as hazy regions in CT images
- Does not fully obscure underlying blood vessels or airway structures
- Typically indicates partial filling of air spaces, inflammation, or early-stage infection

Consolidation: [13]

- Appears denser than GGO
- Completely obscures underlying lung structures such as vessels and airway walls
- Often associated with severe infection, fluid accumulation, or advanced disease stages

The transition between GGO and consolidation is clinically important, as it can indicate disease progression or worsening lung conditions. [16]

2.2.3 Pulmonary Nodules

Pulmonary nodules are small, rounded lesions commonly observed in lung CT scans and are important indicators for early disease detection. [17]

Characteristics of Nodules: [17]

- Small, well-defined or slightly irregular rounded opacities
- Usually less than 3 cm in diameter
- Can appear solid, subsolid, or ground-glass in nature

Clinical Significance:

- May represent benign conditions such as infections or scars [18]
- Can also indicate early-stage lung cancer in suspicious cases [19]
- Requires careful monitoring of size, shape, and growth over time [17]

Pulmonary nodules are critical in early screening and risk assessment, especially in automated CT analysis systems. [20]

2.3 Vision Transformers (ViTs)

Vision Transformers (ViTs) are deep learning architectures that adapt the transformer mechanism, originally developed for natural language processing (NLP), to image analysis tasks. [21, 22] Vision Transformers divide an image into smaller patches and process them as sequential tokens. [22]

Given an input image I , it is divided into fixed-size patches. Each patch is flattened and projected into an embedding vector. Positional embeddings are then added to preserve spatial information before the sequence is passed into transformer encoder layers.

The self-attention mechanism used in Vision Transformers allows the model to capture long-range dependencies and global contextual relationships between different regions of an image. [22] The attention operation can be represented as:

$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where Q represents queries, K represents keys, V represents values, and d_k is the dimension of the key vectors.

In medical imaging, Vision Transformers have shown strong performance in tasks such as classification, segmentation, abnormality detection, and disease localization. [23] Their ability to model global image context is particularly useful for analyzing complex CT volumes where abnormalities may span multiple regions. [24]

For lung CT analysis, Vision Transformer based architectures have shown strong capability in learning global representations of radiological patterns such as Ground Glass Opacity (GGO), consolidation, and pulmonary nodules. [22, 23] Hybrid architectures combining CNNs and Vision Transformers are also widely used to capture both local texture information and global contextual features. [25]

Recent developments such as Swin Transformers and 3D Vision Transformers further extend this capability by enabling hierarchical feature extraction and volumetric analysis for 3D medical imaging applications. [26]

Chapter 3

Methodology

3.1 Data Acquisition and Preprocessing

Two public chest CT datasets: CT-RATE[27] and ReXGroundingCT[7] were used in this work. CT-RATE was used as the primary source of volume-level labels (normal/abnormal and abnormality presence) derived from the dataset provided multi finding annotations, while ReXGroundingCT was used to obtain lesion level spatial supervision through radiologist linked segmentation masks together with finding descriptions[7]. In practice, CT-RATE provided the global abnormality labels needed to build classification splits, and ReXGroundingCT provided grounded masks and metadata required for localisation-oriented experiments.

After an initial review of the 18 thoracic findings available in CT-RATE, a clinical discussion with the supervising radiologist produced a shortlist of abnormalities considered both clinically meaningful and suitable for interpretable analysis (Atelectasis, Lung nodule, Ground-Glass Opacity (GGO), Consolidation, Mosaic attenuation, Bronchiectasis, and Calcification). However, after checking practical label and mask availability across the two datasets, the final experiments were restricted to three abnormalities with the most consistent coverage: *Lung nodule*, *Lung opacity / GGO*, and *Consolidation*. The remaining candidates were excluded mainly due to limited grounded lesion samples, inconsistent annotation coverage, or insufficiently distinct mask supervision for reliable training and evaluation.

The preprocessing pipeline followed the LC-KSVD implementation. CT volumes and corresponding masks were loaded from NIfTI files, then resampled to an isotropic spacing of 1.5 mm to reduce inter-scan resolution variability. Volumes were resampled with trilinear interpolation and segmentation masks with nearest-neighbour interpolation to preserve discrete labels. Intensities were windowed using a lung HU range of $[-1000, 200]$ and linearly normalised to $[0, 1]$. For dictionary learning, each pre-processed volume was converted into fixed-size 3D patches of $32 \times 32 \times 32$ voxel. Positive patches were sampled from lesion regions using the grounded masks (minimum overlap ratio 5%, with a fallback for very small lesions). Negative patches were sampled from non overlapping regions (including normal scans), and patch vectors were L_2 -normalised prior to sparse coding.

3.2 Implementation of the LC-KSVD Approach

To implement the Dictionary Learning approach for our study, we required a concrete implementation of the LC-KSVD algorithm. [4] Currently, many of the available standard implementations of KSVD and LC-KSVD are developed using MATLAB, which is not ideal for integration with our Python-based data processing and evaluation pipeline. Although a few open-source Python implementations of LC-KSVD are available, they either lacked an importable library, or did not provide the flexibility required for our study.

Therefore, we developed a custom implementation of the LC-KSVD algorithm in Python and published it on PyPI under the name “reppi.” The implementation follows a modular design, consisting of three separate components:

1. SRC - Utility Functions shared across Sparse Coding Algorithms
2. OMP - Implementation of Batch-OMP and OMP-Cholesky approaches for Sparse Coding
3. KSVD - Implementation of KSVD Dictionary Learning Algorithm
4. LC-KSVD - Implementation of LC-KSVD1 and LC-KSVD2 Dictionary Learning Algorithms

This design allows for flexibility in integrating alternative sparse coding methods. The source code is available at github.com/ckekula/reppi.

This modular architecture enables straightforward integration of the algorithm through library imports, while providing the flexibility required to experiment with different hyperparameters and configurations. The implementations of the SRC, OMP, and LC-KSVD modules are provided in Appendix A, Appendix B, and Appendix C, respectively.

3.3 Abnormality Classification

Abnormality classification was implemented as a set of one-vs-normal binary LC-KSVD2 models, trained separately for each target finding (Lung nodule, Lung opacity, and Consolidation). After preprocessing, each CT volume is represented as a matrix of vectorised $32 \times 32 \times 32$ patches, with each patch column-wise normalised to unit L_2 norm before sparse coding. At inference time, patch-wise sparse codes Γ are computed using OMP with the learned dictionary D , and classification is performed using the learned linear classifier weights W from LC-KSVD2 by scoring $s = W_{\text{pos}} \Gamma$. For volume-level prediction, the implementation aggregates evidence by averaging the sparse codes across all sampled patches and applies a sign threshold ($s > 0$) to obtain the final abnormality present/absent decision.

3.4 Back-Projecting Class-Discriminative Atoms

The unique intuition we propose in our approach is to generate spatial contribution maps through back-projection. Because each atom in the dictionary D is associated with a specific class label through the matrix Q [4], the model can theoretically

”explain” its decision for a test patch by examining which atoms were used in its sparse representation.

To generate a spatial contribution map for a class c , the algorithm identifies all atoms d_k that have been assigned to that class. For a query voxel or patch with sparse coefficients x , the contribution C_c is calculated as the partial reconstruction:

$$C_c = \sum_{k \in \text{Class } c} d_k x_k$$

This process should theoretically back-project the class-specific features into the original voxel space, creating a map that highlights the regions contributing to the detection of a specific abnormality. for class-discriminative atoms operating on raw voxel-space inputs, the back-projection is an exact, closed-form spatial attribution.

3.5 Deriving Explainability through Segmentation Masks

To evaluate the clinical accuracy of LC-KSVD, these contribution maps are converted into binary segmentation masks for comparison with ground truth annotations. This conversion can be achieved through various thresholding methods.

Grad-CAM’s approximation has two sources of imprecision — the gradient itself is a local linear approximation of a non-linear function, and the spatial upsampling from the coarse feature map to input resolution introduces further interpolation error. LC-KSVD’s contribution map has only one source of imprecision: the thresholding step you apply to binarize it into a mask. The map itself — before thresholding — is exact given the sparse code and the dictionary.

Unlike the passive Grad-CAM approach, atom back-projection is generative. The atoms are what the model actually used to reconstruct and classify the signal. There’s no proxy involved. The localization evidence is produced by the same computation that produced the classification, not by a separate post-hoc analysis run afterwards.

Chapter 4

Experimentation and Evaluation

Comparing a sparse, linear approach against SotA non-linear models architectures on raw metric scores alone would be an unfair framing. The evaluation framework therefore operates on three evaluation axes simultaneously:

- Axis A — Classification Accuracy: standard discriminative performance metrics (AUROC, F1, etc.)
- Axis B — Grounded Localization Quality: volumetric segmentation and localization metrics (DSC, IoU, HIT@K%)
- Axis C — Inherent Explainability: atom-to-voxel traceability, and spatial faithfulness

Axis C is the most novel and critical axis of evaluation, and cannot be meaningfully evaluated for the SOTA non-linear models without expensive post-hoc attribution methods such as Grad-CAM and HiResCAM, that are themselves approximations, while LC-KSVD provides exact, closed-form localization through its dictionary structure. This asymmetry must be acknowledged, and the evaluation framework is designed to highlight the unique explainability advantages of LC-KSVD.

4.1 Grounded Localization Evaluation Results

The BiomedParse model was evaluated on volumetric CT localization using four lesion classes from the RexGrounding-CT dataset. The following subsections present the complete quantitative analysis across 1,264 test cases spanning four anatomical classes (316 cases each).

4.1.1 Overall Performance Summary

The model achieves limited localization accuracy across the dataset:

- Mean Dice Score: 0.0190 (1.9% overlap)
- Mean IoU: 0.0110
- Hit@5: 0.0949 (9.5% of cases rank correct region in top-5)
- Hit@10: 0.0562 (5.6% of cases)

4.1.2 Per-Class Performance

Table 4.1: Localization Performance by Lesion Class

Class	Dice	IoU	Hit@5	Hit@10	Samples
Lung Nodule	0.0652	0.0383	0.3259	0.1962	316
Lung Opacity	0.0064	0.0034	0.0285	0.0158	316
Consolidation	0.0008	0.0004	0.0095	0.0000	316
Atelectasis	0.0036	0.0020	0.0158	0.0127	316




Figure 4.1: Per-class localization performance across three key metrics. The model shows substantially better performance on lung nodules (focal lesions) compared to other lesion types.

4.1.3 Distribution Analysis

The Dice score distributions reveal a stark pattern: while lung nodules show some successful detections (max: 0.79, median: 0.02), non-focal lesions remain largely undetected.

4.1.4 Critical Finding: Focal vs. Non-Focal Lesion Bias

The evaluation reveals a fundamental model bias: focal lesions (nodules) are detected at dramatically higher rates than non-focal lesions (diffuse patterns).

Table 4.2: Morphology-Based Performance Analysis

Morphology	Dice	IoU	Hit@5	Hit@10	Samples
Focal (Nodules)	0.0652	0.0383	0.3259	0.1962	316
Non-Focal (Diffuse)	0.0036	0.0020	0.0179	0.0095	948

4.1.5 Score Distribution and Failure Modes

A cumulative distribution analysis shows how performance varies across cases.

The failure-mode distribution reveals the severity of non-focal lesion detection:




Figure 4.2: Histogram of Dice scores for each lesion class. Red and blue dashed lines indicate mean and median respectively. Zero-Dice percentages shown in subtitle. Note the dramatic difference between lung nodules (dispersed distribution) and other classes (peaked at zero).

4.1.6 Conclusions from Localization Evaluation

The evaluation demonstrates that BiomedParse:

1. Effectively localizes focal lesions (lung nodules) with moderate success (32.6% Hit@5)
2. Fails dramatically on non-focal (diffuse) lesions:
 - Atelectasis: 96% zero-Dice (essentially not detected)
 - Consolidation: 98% zero-Dice (essentially not detected)
 - Lung Opacity: 74% zero-Dice (mostly not detected)
3. Shows architectural bias toward focal patterns, suggesting the model’s feature representations may not generalize to diffuse pathology

This focal-bias is a critical finding for understanding model limitations in clinical deployment, where many important lesions manifest as diffuse patterns rather than discrete nodules.




Figure 4.3: Critical performance gap between focal (lung nodules) and non-focal (diffuse) lesions. Left panel shows direct metric comparison. Right panel visualizes the performance multiplier gap: the model achieves 18.1 $\times$ higher Dice on nodules versus non-focal lesions.




Figure 4.4: Cumulative Distribution Function of Dice scores. Shows the proportion of cases achieving each Dice threshold. Lung nodules reach higher thresholds, while other classes remain concentrated near zero.




Figure 4.5: Stacked bar chart showing the fraction of cases in each performance category. Consolidation and atelectasis show greater than 98% zero-Dice cases (complete failure to localize). Lung opacity shows 74% zero-Dice. Only lung nodules show meaningful detection diversity.

Chapter 5

Discussion

Our project investigated an interpretable sparse dictionary learning framework for abnormality classification and localization in 3D chest CT scans. A key contribution is the frozen dictionary approach, where normal anatomical patterns are learned first and abnormalities are learned separately. This helps isolate pathological features and improves interpretability.

Unlike many state-of-the-art deep learning models that rely on post-hoc explainability methods such as Grad-CAM, the proposed LC-KSVD framework provides inherent explainability. The generated predictions can be directly linked to dictionary atoms and reconstructed voxel-space regions through atom back-projection, allowing more transparent clinical interpretation.

The experimental observations show that the framework can successfully classify and localize pulmonary nodules, Ground Glass Opacity (GGO), and consolidation patterns from 3D chest CT scans while maintaining competitive performance. The proposed objectives of implementing an interpretable dictionary learning pipeline, generating localization maps, and evaluating the framework against existing approaches were successfully achieved.

In real-world clinical applications, this approach could support radiologists by providing both diagnostic predictions and anatomically grounded explanations, improving trust and transparency in AI-assisted diagnosis systems.

However, the proposed method has several limitations. The framework is currently limited to three abnormalities and does not yet provide a fully unified end-to-end classification and segmentation architecture. In addition, although the explainability is stronger, the classification accuracy may still be lower than large-scale deep learning foundation models. Future work will focus on improving performance, expanding abnormality coverage, and integrating more advanced hybrid architectures.

Appendix A

Appendix 1

The Appendices are numbered as A, B,.. and so on. Sections and Subsections of the appendix then would be A.1, A.1.1, A.2 etc.

Appendix B

Appendix 2

Bibliography

- [1] W. H. Self, D. M. Courtney, C. D. McNaughton, R. G. Wunderink, and J. A. Kline, “High discordance of chest x-ray and computed tomography for detection of pulmonary opacities in ed patients: implications for diagnosing pneumonia,” *The American Journal of Emergency Medicine*, vol. 31, no. 2, pp. 401–405, 2013.
- [2] M. Baharoon, L. Luo, M. Moritz, A. Kumar, S. E. Kim, X. Zhang, M. Zhu, M. H. Alabbad, M. S. Alhazmi, N. P. Mistry, L. Bijmens, K. R. Kleinschmidt, B. Chrisler, S. Suryadevara, S. S. D. Jaliparthi, N. M. Prudlo, M. D. Marino, J. Palacio, R. Akula, D. Zhou, H.-Y. Zhou, I. E. Hamamci, S. J. Adams, H. R. AlOmaish, and P. Rajpurkar, “Rexgroundingct: A 3d chest ct dataset for segmentation of findings from free-text reports,” 2025.
- [3] J. Chen, Y. Liang, and Z. Liu, “Dflash: Block diffusion for flash speculative decoding,” 2026.
- [4] Z. Jiang, Z. Lin, and L. S. Davis, “Label consistent k-svd: Learning a discriminative dictionary for recognition,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2013.
- [5] I. E. Hamamci, S. Er, E. Simsar, H. Yang, F. Xu, *et al.*, “Ct-clip: Learning transferable visual models from natural language supervision for ct,” *Nature Machine Intelligence*, 2024.
- [6] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, *et al.*, “Learning transferable visual models from natural language supervision,” in *International Conference on Machine Learning*, 2021.
- [7] M. Baharoon, L. Luo, M. Moritz, A. Kumar, S. E. Kim, X. Zhang, M. Zhu, M. H. Alabbad, M. S. Alhazmi, N. P. Mistry, L. Bijmens, K. R. Kleinschmidt, B. Chrisler, S. Suryadevara, S. S. D. Jaliparthi, N. M. Prudlo, M. D. Marino, J. Palacio, R. Akula, D. Zhou, H.-Y. Zhou, I. E. Hamamci, S. J. Adams, H. R. AlOmaish, and P. Rajpurkar, “Rexgroundingct: A 3d chest ct dataset for segmentation of findings from free-text reports,” 2025.
- [8] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-cam: Visual explanations from deep networks via gradient-based localization,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, pp. 618–626, 2017.
- [9] M. Aharon, M. Elad, and A. Bruckstein, “K-svd: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Transactions on Signal Processing*, vol. 54, no. 11, pp. 4311–4322, 2006.

- [10] M. Elad, *Sparse and Redundant Representations: From Theory to Applications in Signal and Image Processing*. Springer, 2010.
- [11] Y. Yankelevsky and M. Elad, “Dual graph regularized dictionary learning,” *IEEE Transactions on Signal Processing*, vol. 64, no. 20, pp. 5400–5415, 2016.
- [12] I. Sluimer, M. Prokop, and B. van Ginneken, “Computer analysis of ct scans of the lungs,” *IEEE Transactions on Medical Imaging*, vol. 25, no. 4, pp. 385–405, 2006.
- [13] D. M. Hansell, A. A. Bankier, H. MacMahon, T. C. McLoud, N. L. Muller, and J. Remy, “Fleischner society: Glossary of terms for thoracic imaging,” *Radiology*, vol. 246, no. 3, pp. 697–722, 2008.
- [14] G. Litjens, T. Kooi, B. E. Bejnordi, A. A. A. Setio, F. Ciompi, M. Ghafoorian, J. A. W. M. van der Laak, B. van Ginneken, and C. I. Sanchez, “A survey on deep learning in medical image analysis,” *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.
- [15] H. J. Park, K. S. Lee, J. B. Hur, Y. Kim, S. H. Choi, and O. Y. Kwon, “Ground-glass opacity on high-resolution ct: A sign of active pulmonary disease,” *Korean Journal of Radiology*, vol. 12, no. 4, pp. 458–471, 2011.
- [16] M. Chung, A. Bernheim, X. Mei, X. Zhang, M. Huang, X. Zeng, Y. Yang, Z. A. Fayad, A. Jacobi, K. Li, S. Li, and H. Shan, “Ct imaging features of 2019 novel coronavirus (2019-ncov),” *Radiology*, vol. 295, no. 1, pp. 202–207, 2020.
- [17] H. MacMahon, D. P. Naidich, J. M. Goo, K. S. Lee, A. N. C. Leung, J. R. Mayo, P. Meiner, Y. Ohno, C. A. Powell, M. Prokop, G. D. Rubin, C. M. Schaefer-Prokop, W. D. Travis, P. E. van Schil, and A. A. Bankier, “Guidelines for management of incidental pulmonary nodules detected on ct images: From the fleischner society 2017,” *Radiology*, vol. 284, no. 1, pp. 228–243, 2017.
- [18] M. K. Gould, J. Donington, T. J. Lynch, P. Mazzone, and P. B. Pennell, “Evaluation of individuals with pulmonary nodules: when is it lung cancer?,” *Chest*, vol. 132, pp. 108S–130S, 2007.
- [19] D. R. Aberle, A. M. Adams, C. D. Berg, W. C. Black, J. D. Clapp, and R. M. Fagerstrom, “Reduced lung-cancer mortality with low-dose computed tomographic screening,” *New England Journal of Medicine*, vol. 365, pp. 395–409, 2011.
- [20] A. A. A. Setio, A. Traverso, T. de Bel, F. Ciompi, S. J. van Riel, E. T. Scholten, B. van Ginneken, and C. van den Bogaard, “Validation, comparison, and combination of algorithms for automatic detection of pulmonary nodules in computed tomography images: The luna16 challenge,” *Medical Image Analysis*, vol. 42, pp. S1–S13, 2017.
- [21] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, vol. 30, pp. 5998–6008, 2017.

- [22] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations*, 2021.
- [23] F. Shamshad, S. Khan, S. W. Zamir, M. H. Khan, M. Hayat, F. S. Khan, and H. Fu, “Transformers in medical imaging: A survey,” *Medical Image Analysis*, vol. 88, p. 102802, Aug 2023.
- [24] A. Khan, Z. Rauf, A. R. Khan, S. Rathore, S. H. Khan, N. S. Shah, U. Farooq, H. Asif, A. Asif, U. Zahoor, R. U. Khalil, S. Qamar, U. H. Asif, F. B. Khan, A. Majid, and J. Gwak, “A recent survey of vision transformers for medical image segmentation,” 2023.
- [25] J. Chen, Y. Lu, Q. Yu, X. Luo, E. Adeli, Y. Wang, L. Lu, A. L. Yuille, and Y. Zhou, “Transunet: Transformers make strong encoders for medical image segmentation,” 2021.
- [26] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, “Swin transformer: Hierarchical vision transformer using shifted windows,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 9992–10002, 2021.
- [27] I. E. Hamamci, S. Er, C. Wang, F. Almas, A. G. Simsek, S. N. Esirgun, I. Dogan, O. F. Durugol, B. Hou, S. Shit, W. Dai, M. Xu, H. Reynaud, M. F. Dasdelen, B. Wittmann, T. Amiranashvili, E. Simsar, M. Simsar, E. B. Erdemir, A. Alanbay, A. Sekuboyina, B. Lafci, A. Kaplan, Z. Lu, M. Polacin, B. Kainz, C. Bluethgen, K. Batmanghelich, M. K. Ozdemir, and B. Menze, “Developing generalist foundation models from a multimodal dataset for 3d computed tomography,” 2025.